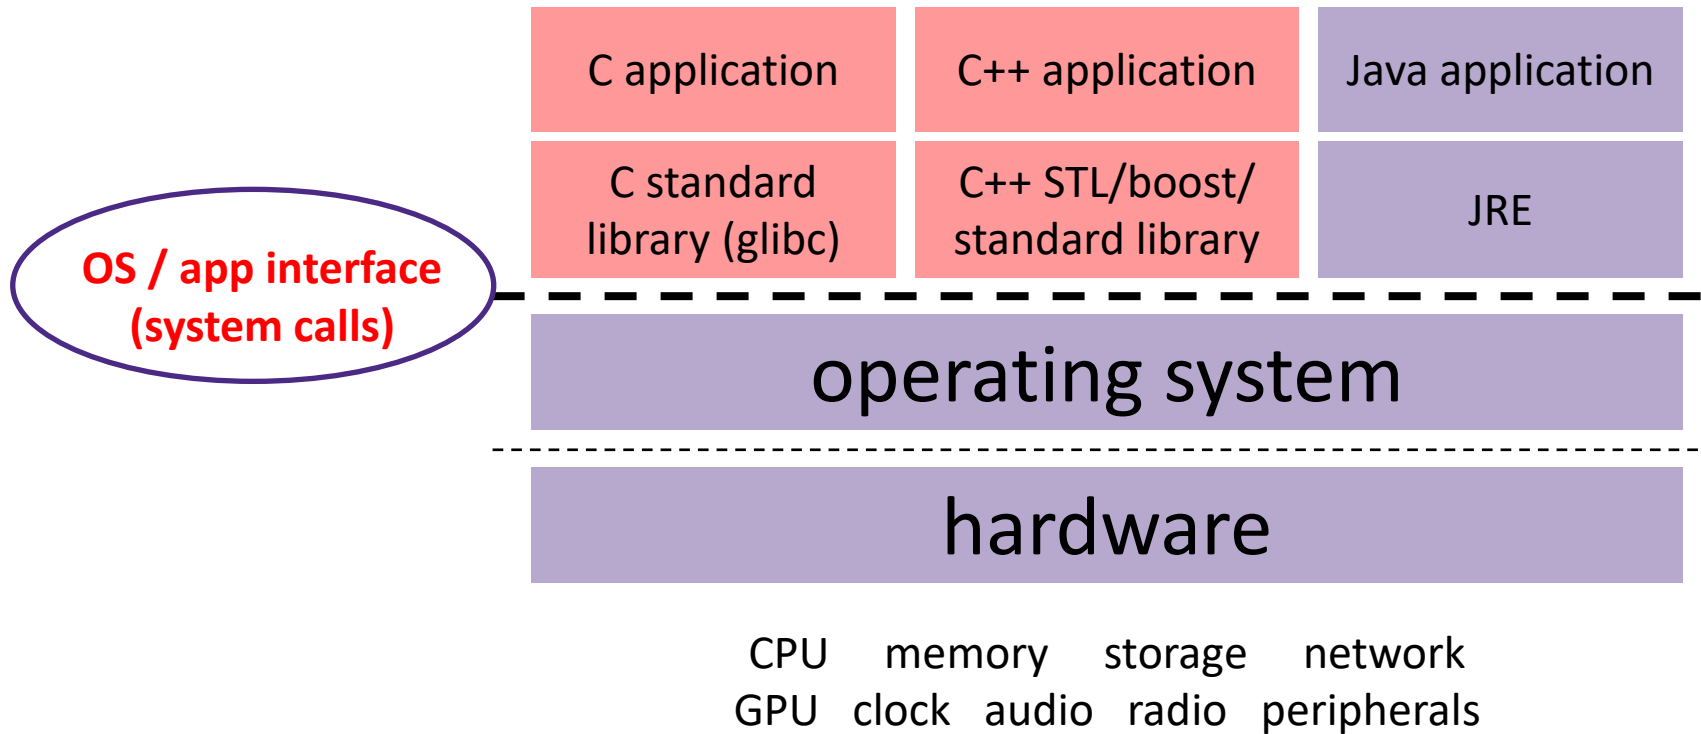


CS/SE 3377

FILE API

Sridhar Alagar

System Calls



From C to POSIX

- ❖ Most Linux versions support a common set of lower-level file access APIs: (conforming to **POSIX** – Portable Operating System Interface)
 - **open()**, **read()**, **write()**, **close()**, **lseek()**
 - Lower-level
 - Also, less convenient

open()/close()

❖ To open a file:

- Pass in the filename and access mode
- Get back a “file descriptor”
 - is just an `int`;
 - Lowest numbered file descriptor available
 - Defaults: `0` is `stdin`, `1` is `stdout`, `2` is `stderr`

```
#include <fcntl.h>    // for open()
#include <unistd.h>    // for close()

...
int fd = open("foo.txt", O_RDONLY);
if (fd == -1) {
    perror("open failed");
    exit(EXIT_FAILURE);
}

...
close(fd);
```

Reading from a file

❖ `ssize_t read(int fd, void* buf, size_t count);`

- Returns the number of bytes read
 - Might be fewer bytes than you requested (!!!)
 - Returns **0** if you're already at the end-of-file
 - Returns **-1** on error (and sets `errno`)
- There are some surprising error modes (check `errno`)
 - **EBADF**: bad file descriptor
 - **EFAULT**: output buffer is not a valid address
 - **EINTR**: read was interrupted, please try again (ARGH!!!!)
 - And many others...

Sample Code

```
int fd1 = open("file.txt"); // returns ?  
read(fd1, buf, 12);  
int fd2 = open("file.txt"); // returns ?  
int fd3 = dup(fd2);          // returns ?  
Close (fd2);
```

fd table

0 ■
1 ■
2 ■
3 ■
4 ■
5 ■

File descriptions(fds)

offset = 0
inode =

inode

location = ...
size = ...

```
int fd1 = open("file.txt"); // returns 3
```

fd table

0 ■
1 ■
2 ■
3 ■
4 ■
5 ■


fds

offset = 12
inode =


inode

location = ...
size = ...


```
int fd1 = open("file.txt"); // returns 3  
read(fd1, buf, 12);
```

```
fd1 = open("file.txt"); // returns 3  
read(fd1, buf, 12);  
fd2 = open("file.txt"); // returns 4
```



```
fd1 = open("file.txt"); // returns 3
read(fd1, buf, 12);
fd2 = open("file.txt"); // returns 4
fd3 = dup(fd2);          // returns 5
```



```
fd1 = open("file.txt"); // returns 3
read(fd1, buf, 12);
fd2 = open("file.txt"); // returns 4
fd3 = dup(fd2);          // returns 5
close(fd2);
```

fd table

0 ■
1 ■
2 ■
3 ■
4 ■
5 ■

fds

offset = 0
inode =

inode

location = ...
size = ...

```
fd1 = open("file.txt"); // returns 3
read(fd1, buf, 12);
fd2 = open("file.txt"); // returns 4
fd3 = dup(fd2);          // returns 5
close(fd2);
close(fd1);
```

Other Low-Level Functions

❖ Read man pages to learn about:

- **write** () – write data
 - `#include <unistd.h>`
- **lseek** () – reposition offset
 - `#include <unistd.h>`
- **fsync** () – flush data to the underlying device
 - `#include <unistd.h>`
- **opendir** (), **readdir** (), **closedir** () – deal with directory listings
 - `#include <dirent.h>`

❖ A useful shortcut sheet (from CMU):

<http://www.cs.cmu.edu/~guna/15-123S11/Lectures/Lecture24.pdf>

One way to read() n bytes

```
int fd = open(filename, O_RDONLY);
char* buf = ...; // buffer of appropriate size
int bytes_left = n;
int result;

while (bytes_left > 0) {
    result = read(fd, buf + (n - bytes_left), bytes_left);
    if (result == -1) {
        if (errno != EINTR) {
            // a real error happened, so return an error result
        }
        // EINTR happened, so do nothing and try again
        continue;
    } else if (result == 0) {
        // EOF reached, so stop reading
        break;
    }
    bytes_left -= result;
}

close(fd);
```

Disclaimer

Some of the materials in this lecture slides are from

- the lecture slides of CS333 Univ of Washington
- the materials prepared by Prof. Arpaci, and Prof. Youjip