

CS/SE 3377

C - Review

Sridhar Alagar

C

❖ Created in 1972 by Dennis Ritchie

- Designed for creating system software
- Portable across machine architectures
- Most recently updated in 1999 (C99) and 2011 (C11)

❖ Characteristics

- “Low-level” language that allows us to exploit underlying features of the architecture – **but easy to fail spectacularly (!)**
- Procedural (not object-oriented)
- “Weakly-typed” or “type-unsafe”
- Small, basic library compared to Java, C++

Generic C Program Layout

```
#include <system_files>
#include "local_files"

#define macro_name macro_expr

/* declare functions */
/* declare external variables & structs */

int main(int argc, char* argv[]) {
    /* the innards */
}

/* define other functions */
```

C: main

- ❖ To get command-line arguments in `main`, use:

```
int main(int argc, char* argv[])
```

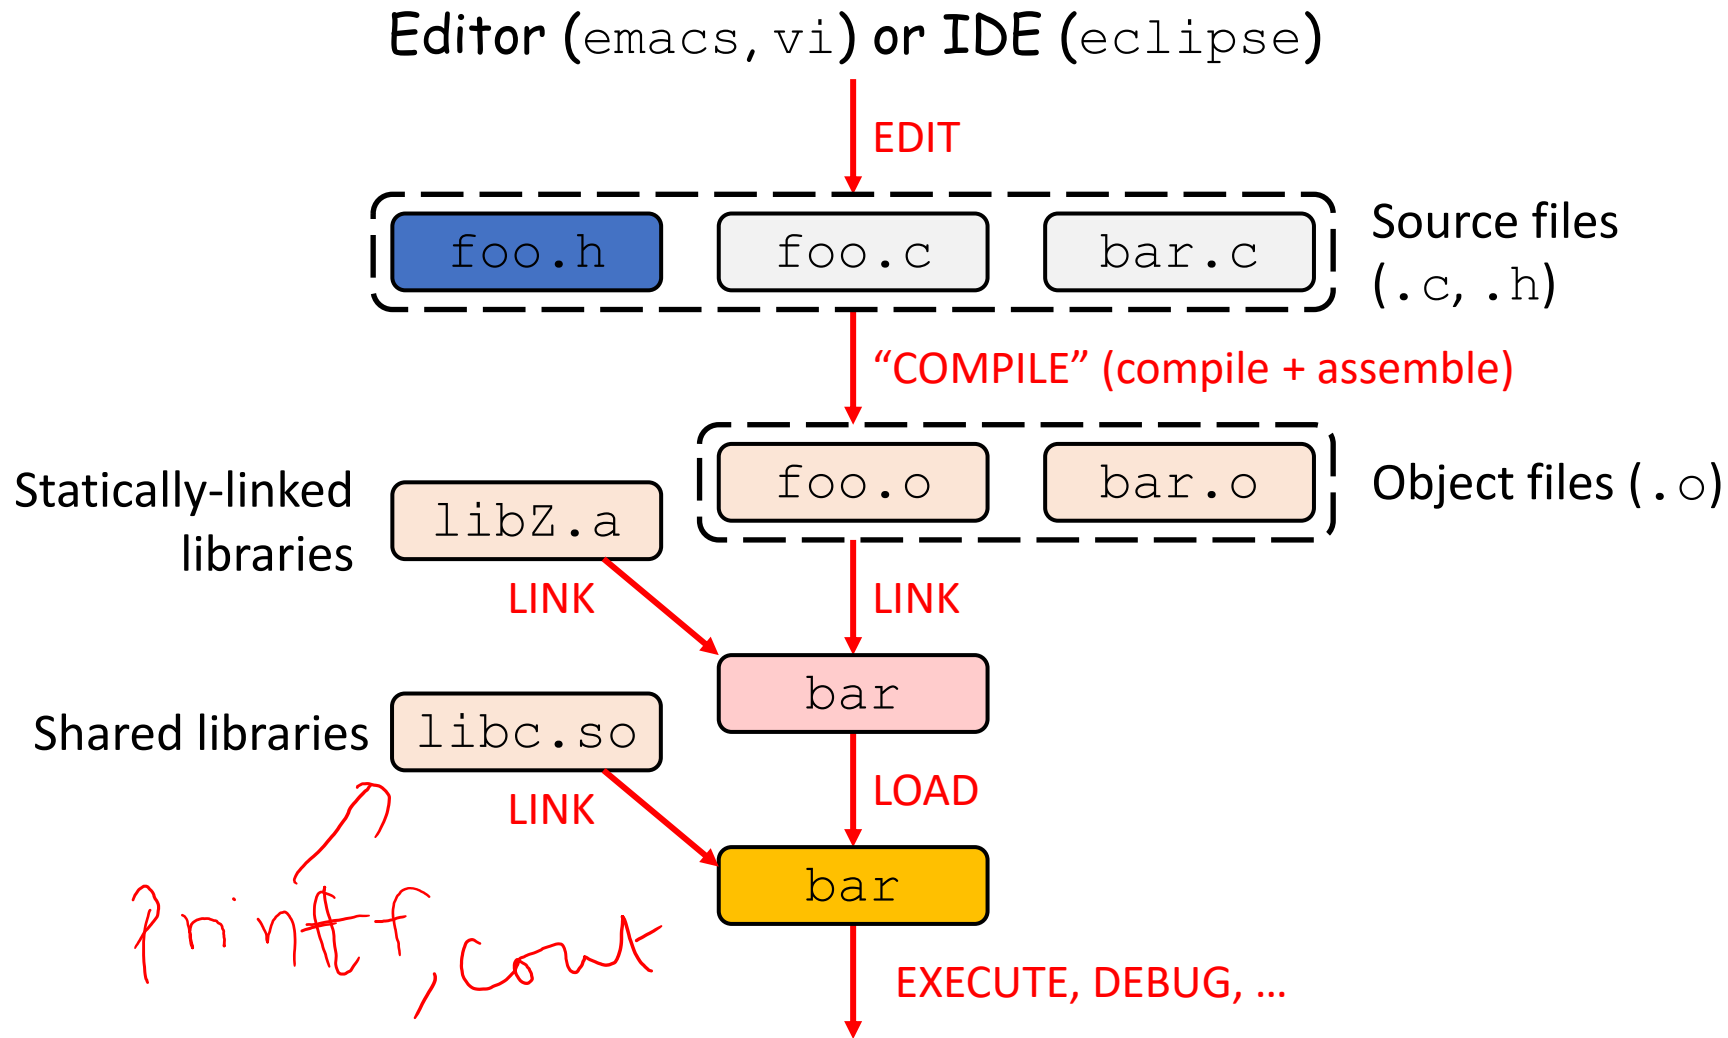
- ❖ What does this mean?

- `argc` contains the number of strings on the command line (the executable name counts as one, plus one for each argument).
- `argv` is an array containing *pointers* to the arguments as strings (more on pointers later)

- ❖ Example: `$ foo hello 877`

- `argc = 3`
- `argv[0] = "foo", argv[1] = "hello", argv[2] = "877"`

C programming to execution



[Check this link for how to create static and shared libraries](#)

C to machine code

```
void sumstore(int x, int y,  
               int* dest) {  
    *dest = x + y;  
}
```

C source file
(sumstore.c)

C compiler (gcc -S)

sumstore:

```
    addl    %edi, %esi  
    movl    %esi, (%rdx)  
    ret
```

Assembly file
(sumstore.s)

Assembler (gcc -c or as)

```
400575: 01 fe  
      89 32  
      c3
```

Machine code
(sumstore.o)

C compiler
(gcc -c)

When things go wrong...

❖ Errors and Exceptions

- C does not have exception handling (no `try/catch`)
- Errors are returned as integer error codes from functions
 - Standard codes found in `stdlib.h`:
`EXIT_SUCCESS` (usually 0) and `EXIT_FAILURE` (non-zero)
 - Return value from `main` is a status code
- Because of this, error handling is ugly and inelegant

❖ Crashes

- If you do something bad, you hope to get a “segmentation fault” (believe it or not, this is the “good” option)

Java vs C

- ❖ Are Java and C mostly similar (S) or significantly different (D) in the following categories?

Language Feature	S/D	Differences in C
Control structures <i>if, for, while</i>	<i>S</i>	
Primitive datatypes <i>char, int</i>	<i>D</i>	<i>char is 1 byte; ASCII</i>
Operators	<i>S</i>	<i>>>> not in C</i>
Casting	<i>D</i>	<i>not strongly type</i>
Arrays	<i>D</i>	<i>not objects.</i>
Memory management	<i>D</i>	<i>malloc(); no garbage collection</i>

Primitive types in C

❖ Integer types

- `char`, `int`

❖ Floating point

- `float`, `double`

❖ Modifiers

- `short` [int]
- `long` [int, double]
- `signed` [char, int]
- `unsigned` [char, int]

C Data Type	32-bit	64-bit	printf
char	1	1	%c
short int	2	2	%hd
unsigned short int	2	2	%hu
int	4	4	%d / %i
unsigned int	4	4	%u
long int	4	8	%ld
long long int	8	8	%lld
float	4	4	%f
double	8	8	%lf
long double	12	16	%Lf
pointer	4	8	%p

C99 extended integer types

- ❖ Solves the conundrum of “how big is a `long int`?”

```
#include <stdint.h>

void foo(void) {
    int8_t  a;  // exactly 8 bits, signed
    int16_t b;  // exactly 16 bits, signed
    int32_t c;  // exactly 32 bits, signed
    int64_t d;  // exactly 64 bits, signed
    uint8_t w;  // exactly 8 bits, unsigned
    ...
}
```

```
void sumstore(int x, int y, int* dest) {
```

```
void sumstore(int32_t x, int32_t y, int32_t* dest) {
```

Pointers

Pointers - basics

❖ Variables that store addresses

- It points to somewhere in the process' virtual address space
- `&foo` produces the virtual address of `foo`

❖ Generic definition: `type* name;` or `type *name;`

- Recommended: do not define multiple pointers on same line:

`int* p1, p2;` not the same as `int *p1, *p2;`


- Instead, use:
`int *p1;`
`int *p2;`

❖ *Dereference* a pointer using the unary `*` operator


- Access the memory referred to by a pointer

val X
X = 1)

foo = &X



~~foo~~



Pointer example

pointer.c


```
#include <stdio.h>
#include <stdlib.h>
#include <stdint.h>
#include <inttypes.h>

int main(int argc, char** argv) {
    int32_t x = 351;
    int32_t* p;    // p is a pointer to a int

    p = &x;        // p now contains the addr of x
    printf("&x is %p\n", &x);
    printf(" p is %p\n", p);
    printf(" x is %d\n", x);

    *p = 333;      // change the value of x
    printf(" x is %d \n", x);

    return EXIT_SUCCESS;
}
```



~~x~~ p = 333

X

Pointers illustrated

boxarrow.c

```
int main(int argc, char** argv) {
    int x = 1;
    int arr[3] = {2, 3, 4};
    int* p = &arr[1];

    printf("&x: %p;  x: %d\n", &x, x);
    printf("&arr[0]: %p;  arr[0]: %d\n", &arr[0], arr[0]);
    printf("&arr[2]: %p;  arr[2]: %d\n", &arr[2], arr[2]);
    printf("&p: %p; p: %p; *p: %d\n", &p, p, *p);

    return EXIT_SUCCESS;
}
```

address

name	value
------	-------

Pointers illustrated

boxarrow.c

```
int main(int argc, char** argv) {
    int x = 1;
    int arr[3] = {2, 3, 4};
    int* p = &arr[1];

    printf("&x: %p;  x: %d\n", &x, x);
    printf("&arr[0]: %p;  arr[0]: %d\n", &arr[0], arr[0]);
    printf("&arr[2]: %p;  arr[2]: %d\n", &arr[2], arr[2]);
    printf("&p: %p;  p: %p;  *p: %d\n", &p, p, *p);

    return EXIT_SUCCESS;
}
```

address	name	value
---------	------	-------

&x	x	value
&arr[2]	arr[2]	value
&arr[1]	arr[1]	value
&arr[0]	arr[0]	value
&p	p	value

Pointers illustrated

boxarrow.c

```
int main(int argc, char** argv) {
    int x = 1;
    int arr[3] = {2, 3, 4};
    int* p = &arr[1];

    printf("&x: %p;  x: %d\n", &x, x);
    printf("&arr[0]: %p;  arr[0]: %d\n", &arr[0], arr[0]);
    printf("&arr[2]: %p;  arr[2]: %d\n", &arr[2], arr[2]);
    printf("&p: %p;  p: %p;  *p: %d\n", &p, p, *p);

    return EXIT_SUCCESS;
}
```

address	name	value
---------	------	-------


&x	x	1
&arr[2]	arr[2]	4
&arr[1]	arr[1]	3
&arr[0]	arr[0]	2
&p	p	&arr[1]

Pointers illustrated

$p = \text{malloc}(20)$

boxarrow.c

```
int main(int argc, char** argv) {  
    int x = 1;  
    int arr[3] = {2, 3, 4};  
    int* p = &arr[1];  
  
    printf("&x: %p; x: %d\n", &x, x);  
    printf("&arr[0]: %p; arr[0]: %d\n", &arr[0], arr[0]);  
    printf("&arr[2]: %p; arr[2]: %d\n", &arr[2], arr[2]);  
    printf("&p: %p; p: %p; *p: %d\n", &p, p, *p);  
  
    return EXIT_SUCCESS;  
}
```




address	name	value
---------	------	-------

0x7fff...4c	x	1
0x7fff...48	arr[2]	4
0x7fff...44	arr[1]	3
0x7fff...40	arr[0]	2
0x7fff...38	p	0x7fff...44

Pointers key takeaway

“Pointers are just variables that contain memory addresses”

“Since pointers are variables, we can do all these things recursively!”



Pointer arithmetic

- ❖ Pointers are typed
 - Tells the compiler the size of the data you are pointing to
- ❖ Pointer arithmetic is scaled by `sizeof(*p)`
 - Works nicely for arrays
- ❖ Valid pointer arithmetic:
 - Add/subtract an integer to/from a pointer
 - Subtract two pointers (within stack frame or malloc block)
 - Compare pointers (<, <=, ==, !=, >, >=), including NULL
 - ... but plenty of valid-but-inadvisable operations, too

Pointer Arithmetic - Example

```
int main(int argc, char** argv) {  
    int arr[3] = {2, 3, 4};  
    int* p = &arr[0];  
    int* p1 = &arr[1];  
    int* p2 = p1 + 1;  
}
```


$p = \&arr$
 $\&arr = p;$

Running Program's Layout

Compiling Multi-file Programs

❖ The **linker** combines multiple object files plus statically-linked libraries to produce an executable

- Includes many standard libraries (e.g. libc, crt1)
 - A **library** is just a pre-assembled collection of .o files



Handwritten code snippet:


```
sumnum.c
main() {
    sumstore();
    // ...
}
sumstore.c {
    // ...
}
```

Annotations: A red arrow points from the `sumstore()` call in `main()` to the `sumstore.c` file. Another red arrow points from the `sumstore.c` file to the `sumnum` executable.

Handwritten notes:


- return
- fff
- STATIC DATA
- sumnum.c
- CONC

Executing a Program



Process and its memory

- ❖ A process is a program in execution
- ❖ The OS gives each process the illusion of its own private memory
 - Called the process' **address space**
 - 2^{64} bytes on a 64-bit machine

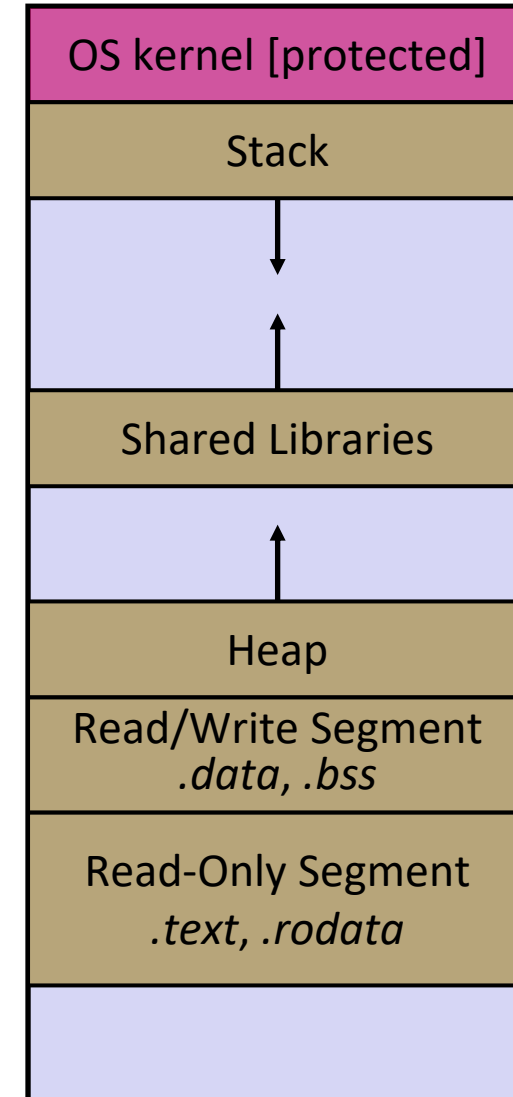


Loading



- ❖ When the OS loads a program it:
 - 1) Creates an address space
 - 2) Inspects the executable file to see what's in it
 - 3) Copies regions of the file into the right place in the address space
 - 4) Does any final linking, relocation, or other needed preparation


0xFF...FF



0x00...00

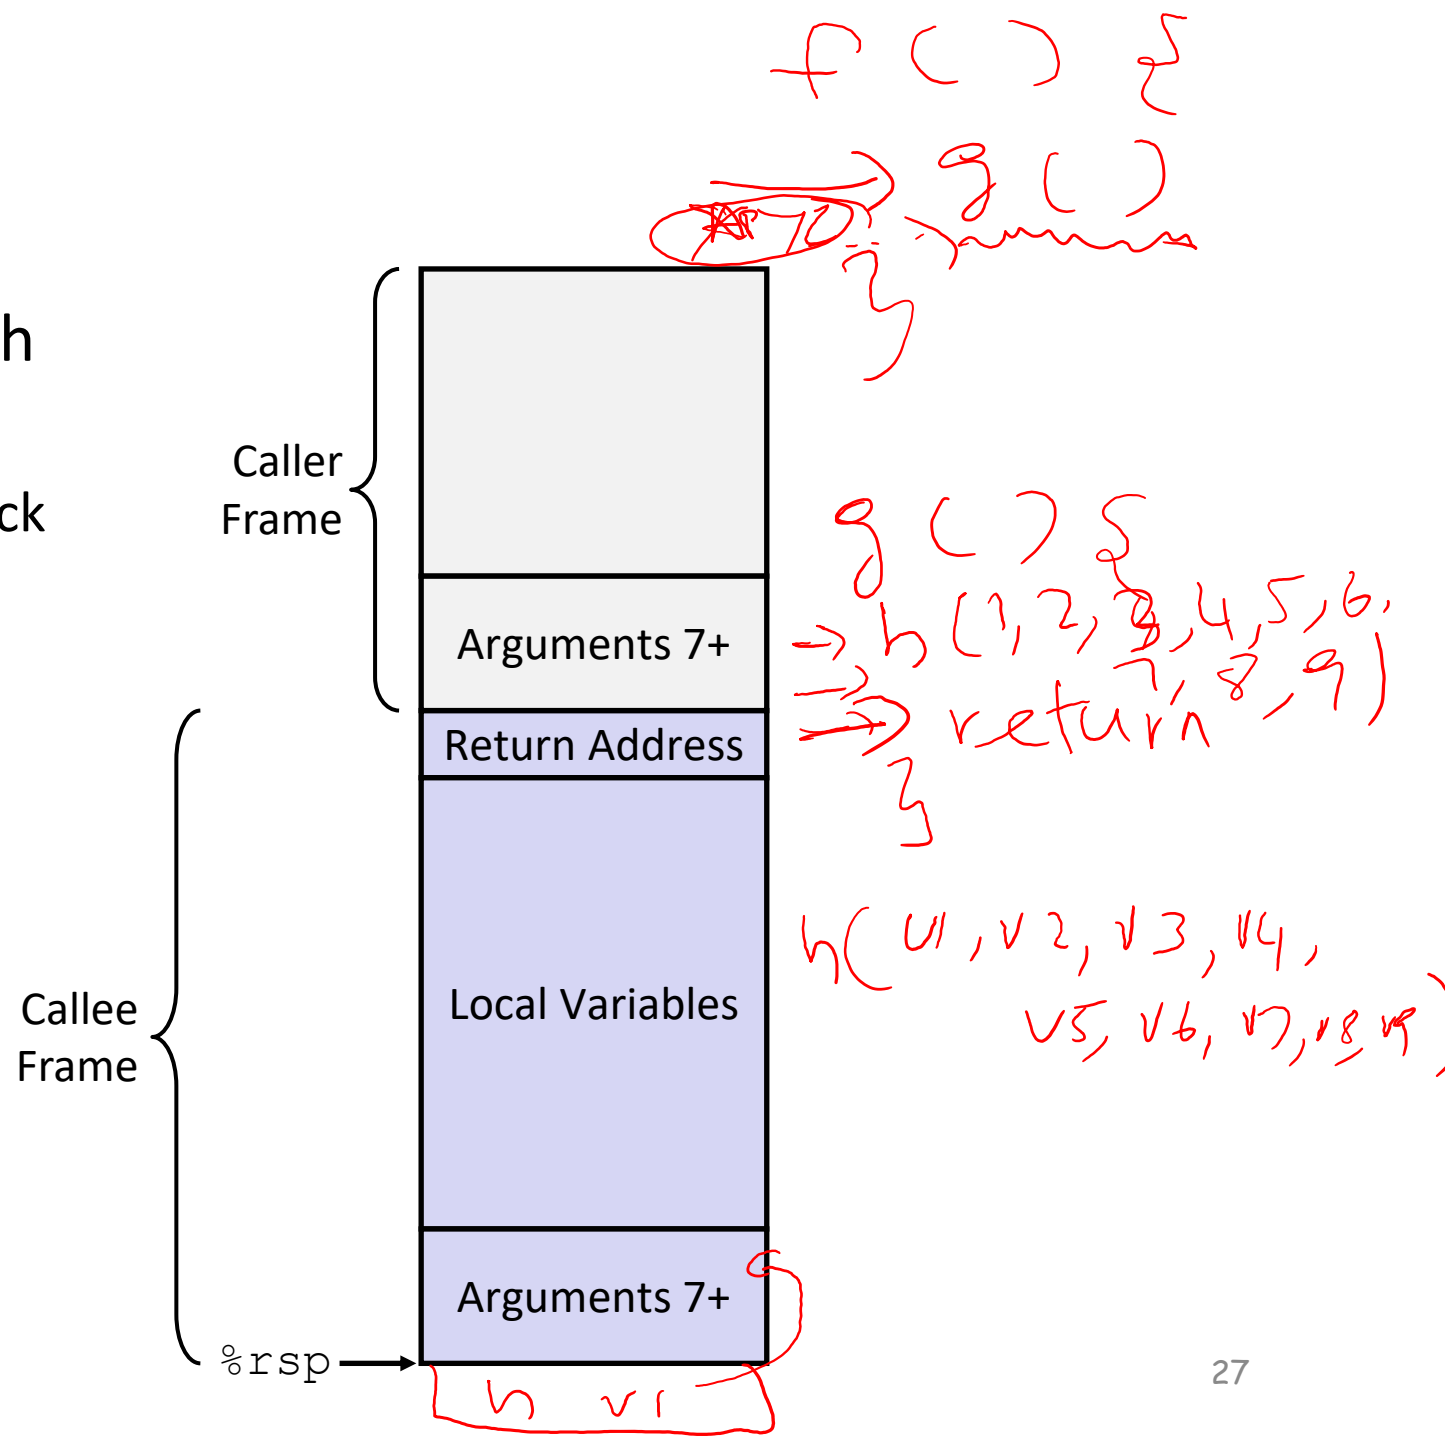
Memory management

- ❖ *Local* variables on the Stack
 - Allocated/freed during functions call/ret (`push`, `pop`, `mov`)
- ❖ *Global* and *static* variables in Data
 - Allocated/freed when the process starts/exits
- ❖ *Dynamically-allocated* data on the Heap
 - `malloc()` to request; `free()` to free, otherwise **memory leak**



The Stack

- ❖ Used to store data associated with function calls
 - Compiler-inserted code manages stack frames for you
- ❖ Stack frame typically includes:
 - Address to return to
 - Local variables
 - Argument build
 - Only if > 6 used



Stack in action

stack.c

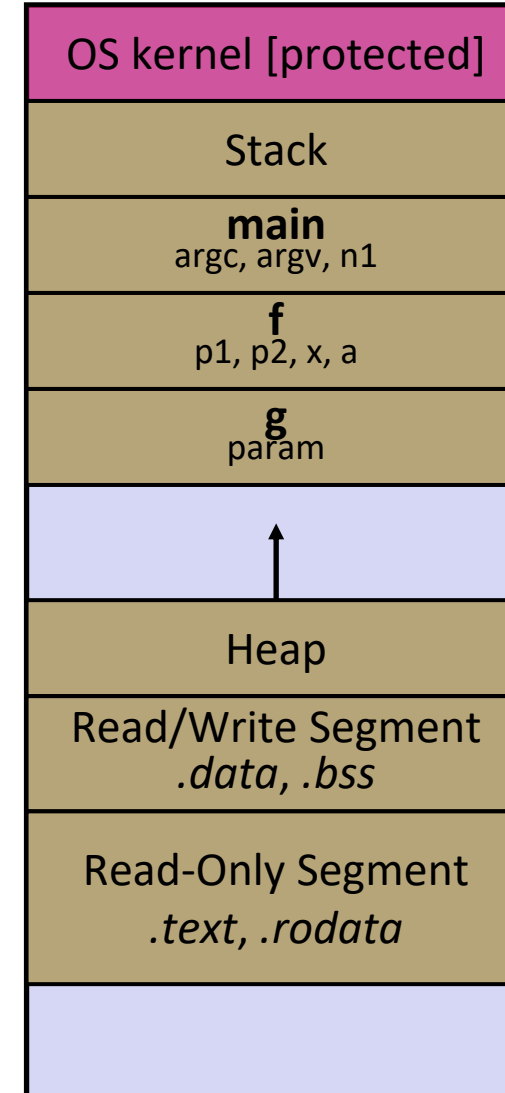
```
int32_t f(int32_t, int32_t);
int32_t g(int32_t);

→ int main(int argc, char** argv) {
→   int32_t n1 = f(3, -5);
→   n1 = g(n1);
→   return EXIT_SUCCESS;
→ }

→ int32_t f(int32_t p1, int32_t p2) {
→   int32_t x;
→   int32_t a[3];
→   ...
→   x = g(a[2]);
→   return x;
→ }

→ int32_t g(int32_t param) {
→   return param * 2;
→ }
```

x = p1, p1 = 2



Stack in action

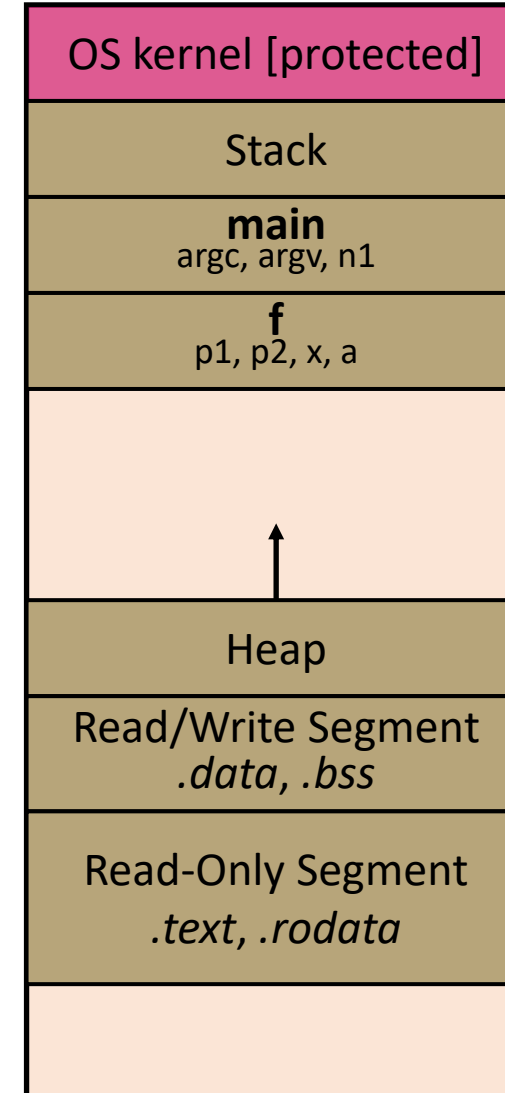
stack.c

```
int32_t f(int32_t, int32_t);
int32_t g(int32_t);

int main(int argc, char** argv) {
    int32_t n1 = f(3, -5);
    n1 = g(n1);
    return EXIT_SUCCESS;
}

int32_t f(int32_t p1, int32_t p2) {
    int32_t x;
    int32_t a[3];
    ...
    x = g(a[2]);
    return x;
}

int32_t g(int32_t param) {
    return param * 2;
}
```



Stack in action

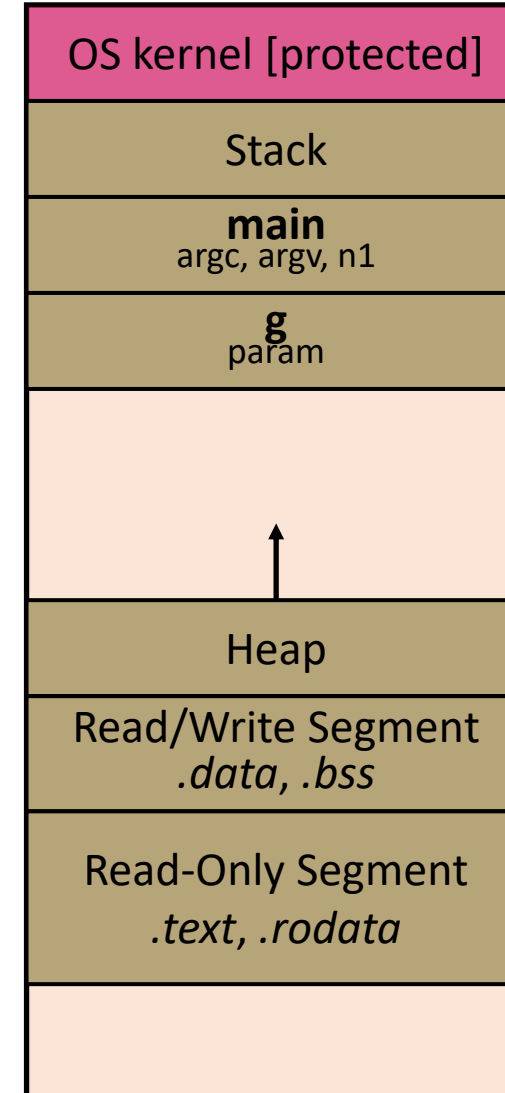
stack.c

```
int32_t f(int32_t, int32_t);
int32_t g(int32_t);

int main(int argc, char** argv) {
    int32_t n1 = f(3, -5);
    n1 = g(n1);
    return EXIT_SUCCESS;
}

int32_t f(int32_t p1, int32_t p2) {
    int32_t x;
    int32_t a[3];
    ...
    x = g(a[2]);
    return x;
}

int32_t g(int32_t param) {
    return param * 2;
}
```



C is Call-By-Value

- ❖ C (and Java) pass arguments by *value*
 - Callee receives a **local copy** of the argument
 - Register or Stack
 - If the callee modifies a parameter, the caller's copy *isn't* modified

```
void swap(int a, int b) {  
    int tmp = a;  
    a = b;  
    b = tmp;  
}  
  
int main(int argc, char** argv) {  
    int a = 42, b = -7;  
    swap(a, b);  
    ...  
}
```

Faking Call-By-Reference in C

- ❖ Can use pointers to *approximate* call-by-reference
 - Callee still receives a **copy** of the pointer (*i.e.* call-by-value), but it can modify something in the caller's scope by dereferencing the pointer parameter

```
void swap(int* a, int* b) {  
    int tmp = *a;  
    *a = *b;  
    *b = tmp;  
}  
  
int main(int argc, char** argv) {  
    int a = 42, b = -7;  
    swap(&a, &b);  
    ...  
}
```


Basic data structure

- ❖ C does not support objects!!!
- ❖ **Arrays** are contiguous chunks of memory
 - Arrays have no methods and do not know their own length
 - Can easily run off ends of arrays in C – **security bugs!!!**
- ❖ **Strings** are null-terminated char arrays
 - Strings have no methods, but **string.h** has helpful utilities

<code>char* x = "hello\n";</code>	<code>x</code>	\rightarrow	<table border="1"><tr><td>h</td><td>e</td><td>l</td><td>l</td><td>o</td><td>\n</td><td>\0</td></tr></table>	h	e	l	l	o	\n	\0
h	e	l	l	o	\n	\0				

- ❖ **Structs** are the most object-like feature, but are just collections of fields – no “methods” or functions

Arrays

- ❖ Definition: `type name [size]`
 - Allocates `size * sizeof (type)` bytes of *contiguous* memory
 - Normal usage is a compile-time constant for `size`
(e.g. `int32_t scores [175] ;`)
- ❖ Size of an array
 - Not stored anywhere – array does not know its own size!
 - `sizeof (array)` only works in variable scope of array definition

Array as parameters

- ❖ It's tricky to use arrays as parameters
 - What happens when you use an array name as an argument?
 - Recall: arrays do not know their own size

```
// prototype
int32_t sumAll(int32_t a[]);

int main(int argc, char** argv) {
    int32_t numbers[] = {9, 8, 1, 9, 5};
    int32_t sum = sumAll(numbers);
    return EXIT_SUCCESS;
}

int32_t sumAll(int32_t a[]) {
    int32_t i, sum = 0;
    for (i = 0; i < ...???)
}
```

Solution: Pass Size as Parameter

```
// prototype
int32_t sumAll(int32_t a[], int size);

int main(int argc, char** argv) {
    int32_t numbers[] = {9, 8, 1, 9, 5};
    int32_t sum = sumAll(numbers, 5);
    printf("sum is: %d\n", sum);
    return EXIT_SUCCESS;
}

int32_t sumAll(int32_t a[], int size) {
    int32_t i, sum = 0;
    for (i = 0; i < size; i++) {
        sum += a[i];
    }
    return sum;
}
```

arraysum.c

- This is the standard idiom in C programs

Returning an Array

- ❖ Local variables, including arrays, are allocated on the Stack
 - They “disappear” when a function returns!
 - Can’t safely return local arrays from functions

```
int32_t* copyArray(int32_t src[], int32_t size) {  
    int32_t i, dst[size];    // OK in C99  
  
    for (i = 0; i < size; i++) {  
        dst[i] = src[i];  
    }  
  
    return dst;    // no compiler error, but wrong!  
}
```

Solution: Output Parameter

- ❖ Create the “returned” array in the caller
 - Pass it as an **output parameter** to `copyarray()`
 - A pointer parameter that allows the called function to store values that the caller can use
 - Works because arrays are “passed” as pointers
 - “Feels” like call-by-reference, *but technically it's not*

```
void copyArray(int32_t src[], int32_t dst[], int32_t size) {  
    int32_t i;  
  
    for (i = 0; i < size; i++) {  
        dst[i] = src[i];  
    }  
}
```

Arrays: Call-By-Value or Call-By-Reference?

- ❖ **Technical answer:** a $T[]$ array parameter is “decayed” to a pointer of type T^* , and the *pointer* is passed by value
 - So it acts like a call-by-reference array (if callee changes the array parameter *elements* it changes the caller’s array)
 - But it’s really a call-by-value pointer (the callee can change the pointer *parameter* to point to something else(!))

```
void copyArray(int32_t src[], int32_t dst[], int32_t size) {  
    int32_t i;  
    int32_t copy[size]; // OK in C99, still stylistically bad  
    for (i = 0; i < size; i++) {  
        copy[i] = src[i];  
    }  
    dst = copy; // doesn't change caller's array  
}
```

Dynamic Allocation

- ❖ What we want is *dynamically*-allocated memory
 - Your program explicitly requests a new block of memory
 - The language allocates it at runtime, perhaps with help from OS
 - Dynamically-allocated memory persists until either:
 - Your code explicitly deallocated it (manual memory management)
 - A garbage collector collects it (automatic memory management)
- ❖ C requires you to manually manage memory
 - Gives you more control, but causes headaches

malloc()

- ❖ General usage: `var = (type*) malloc(size in bytes)`
- ❖ **malloc** allocates a block of memory of the requested size
 - Returns a pointer to the first byte of that memory
 - And **returns NULL** if the memory allocation failed!
 - You should assume that the memory initially contains garbage
 - You'll typically use **sizeof** to calculate the size you need

```
// allocate a 10-float array
float* arr = (float*) malloc(10*sizeof(float));
if (arr == NULL) {
    return errcode;
}
...    // do stuff with arr
```

free()

- ❖ Usage: `free(pointer);`
- ❖ Deallocates the memory pointed-to by the pointer
 - Pointer *must* point to the first byte of heap-allocated memory (*i.e.* something previously returned by `malloc`)
 - Freed memory becomes eligible for future allocation
 - Pointer is unaffected by call to free
 - Defensive programming: can set pointer to `NULL` after freeing it

```
float* arr = (float*) malloc(10*sizeof(float));  
if (arr == NULL)  
    return errcode;  
...           // do stuff with arr  
free(arr);  
arr = NULL;   // OPTIONAL
```

Structured Data

❖ A `struct` is a C datatype that contains a set of fields

- Similar to a Java class, but with no methods or constructors
- Useful for defining new structured types of data
- Behave similarly to primitive variables

❖ Generic declaration:

```
struct tagname {  
    type1 name1;  
    ...  
    typeN nameN;  
};
```

```
// the following defines a new  
// structured datatype called  
// a "struct Point"  
struct Point {  
    float x, y;  
};  
  
// declare and initialize a  
// struct Point variable  
struct Point origin = {0.0, 0.0};
```

Using struct

- ❖ Use “.” to refer to a field in a struct
- ❖ Use “->” to refer to a field from a struct pointer
 - Dereferences pointer first, then accesses field

```
struct Point {  
    float x, y;  
};  
  
int main(int argc, char** argv) {  
    struct Point p1 = {0.0, 0.0}; // p1 is stack allocated  
    struct Point* p1_ptr = &p1;  
  
    p1.x = 1.0;  
    p1_ptr->y = 2.0; // equivalent to (*p1_ptr).y = 2.0;  
    return EXIT_SUCCESS;  
}
```

simplestruct.c

Dynamically allocated Structs

- ❖ You can **malloc** and **free** structs, just like other data type
 - **sizeof** is particularly helpful here

```
// a complex number is a + bi
typedef struct complex_st {
    double real;    // real component
    double imag;    // imaginary component
} Complex, *ComplexPtr;

// note that ComplexPtr is equivalent to Complex*
ComplexPtr AllocComplex(double real, double imag) {
    Complex* retval = (Complex*) malloc(sizeof(Complex));
    if (retval != NULL) {
        retval->real = real;
        retval->imag = imag;
    }
    return retval;
}
```

complexstruct.c

`gdb` - Gnu debugger

- Must learn to use it. Otherwise, debugging can be miserable
- Very useful to understand code also
- Source code should be compiled with '-g' option to use `gdb`
- Check <https://sourceware.org/gdb>

Disclaimer

Some of the materials in this lecture slides are from the lecture slides of CS333 Univ of Washington

Function declaration

- ❖ Informs the compiler arguments and return types; function definitions can then be in a logical order
 - Function comment usually by the *prototype*

```
// sum of integers from 1 to max
int32_t sumTo(int32_t); // func prototype

int main(int argc, char** argv) {
    printf("sumTo(5) is: %d\n", sumTo(5));
    return EXIT_SUCCESS;
}

int32_t sumTo(int32_t max) {
    int32_t i, sum = 0;
    for (i = 1; i <= max; i++) {
        sum += i;
    }
    return sum;
}
```


Function Declaration vs. Definition

- ❖ C/C++ make a careful distinction between these two
- ❖ **Definition:** the *thing* itself
 - *e.g.* code for function, variable definition that creates storage
 - Must be **exactly one** definition of each thing (no duplicates)
- ❖ **Declaration:** description of a thing
 - *e.g.* function prototype, external variable declaration
 - Often in header files and incorporated via `#include`
 - Should also `#include` declaration in the file with the actual definition to check for consistency
 - Needs to appear in **all files** that use that thing
 - Should appear before first use

Multiple C programs

C source file 1
(sumstore.c)

```
void sumstore(int x, int y, int* dest) {  
    *dest = x + y;  
}
```

C source file 2
(sumnum.c)

```
#include <stdio.h>  
  
void sumstore(int x, int y, int* dest);  
  
int main(int argc, char** argv) {  
    int z, x = 351, y = 333;  
    sumstore(x, y, &z);  
    printf("%d + %d = %d\n", x, y, z);  
    return 0;  
}
```

Compile together:

```
$ gcc -o sumnum sumnum.c sumstore.c
```